

Microsoft's DevOps Journey

Paul Hacker

DevOps Architect – DevOps CAT
pahacker@microsoft.com



Introductions

Paul Hacker

Came to Microsoft in November 2017

DevOps Architect - DevOps Customer Advisory Team (CAT)

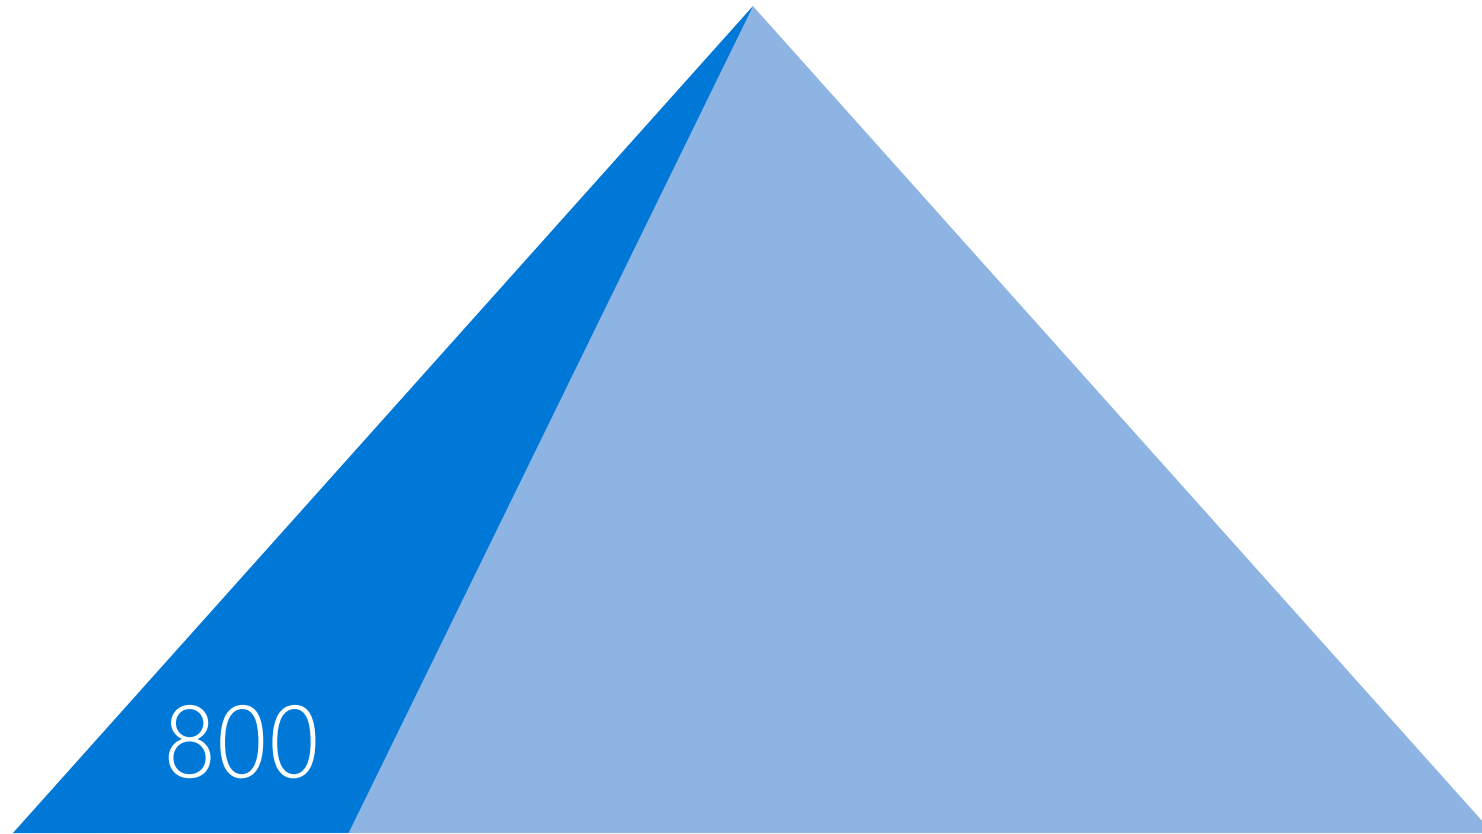
Microsoft MVP in Dev Tools 14 years

Conference speaker



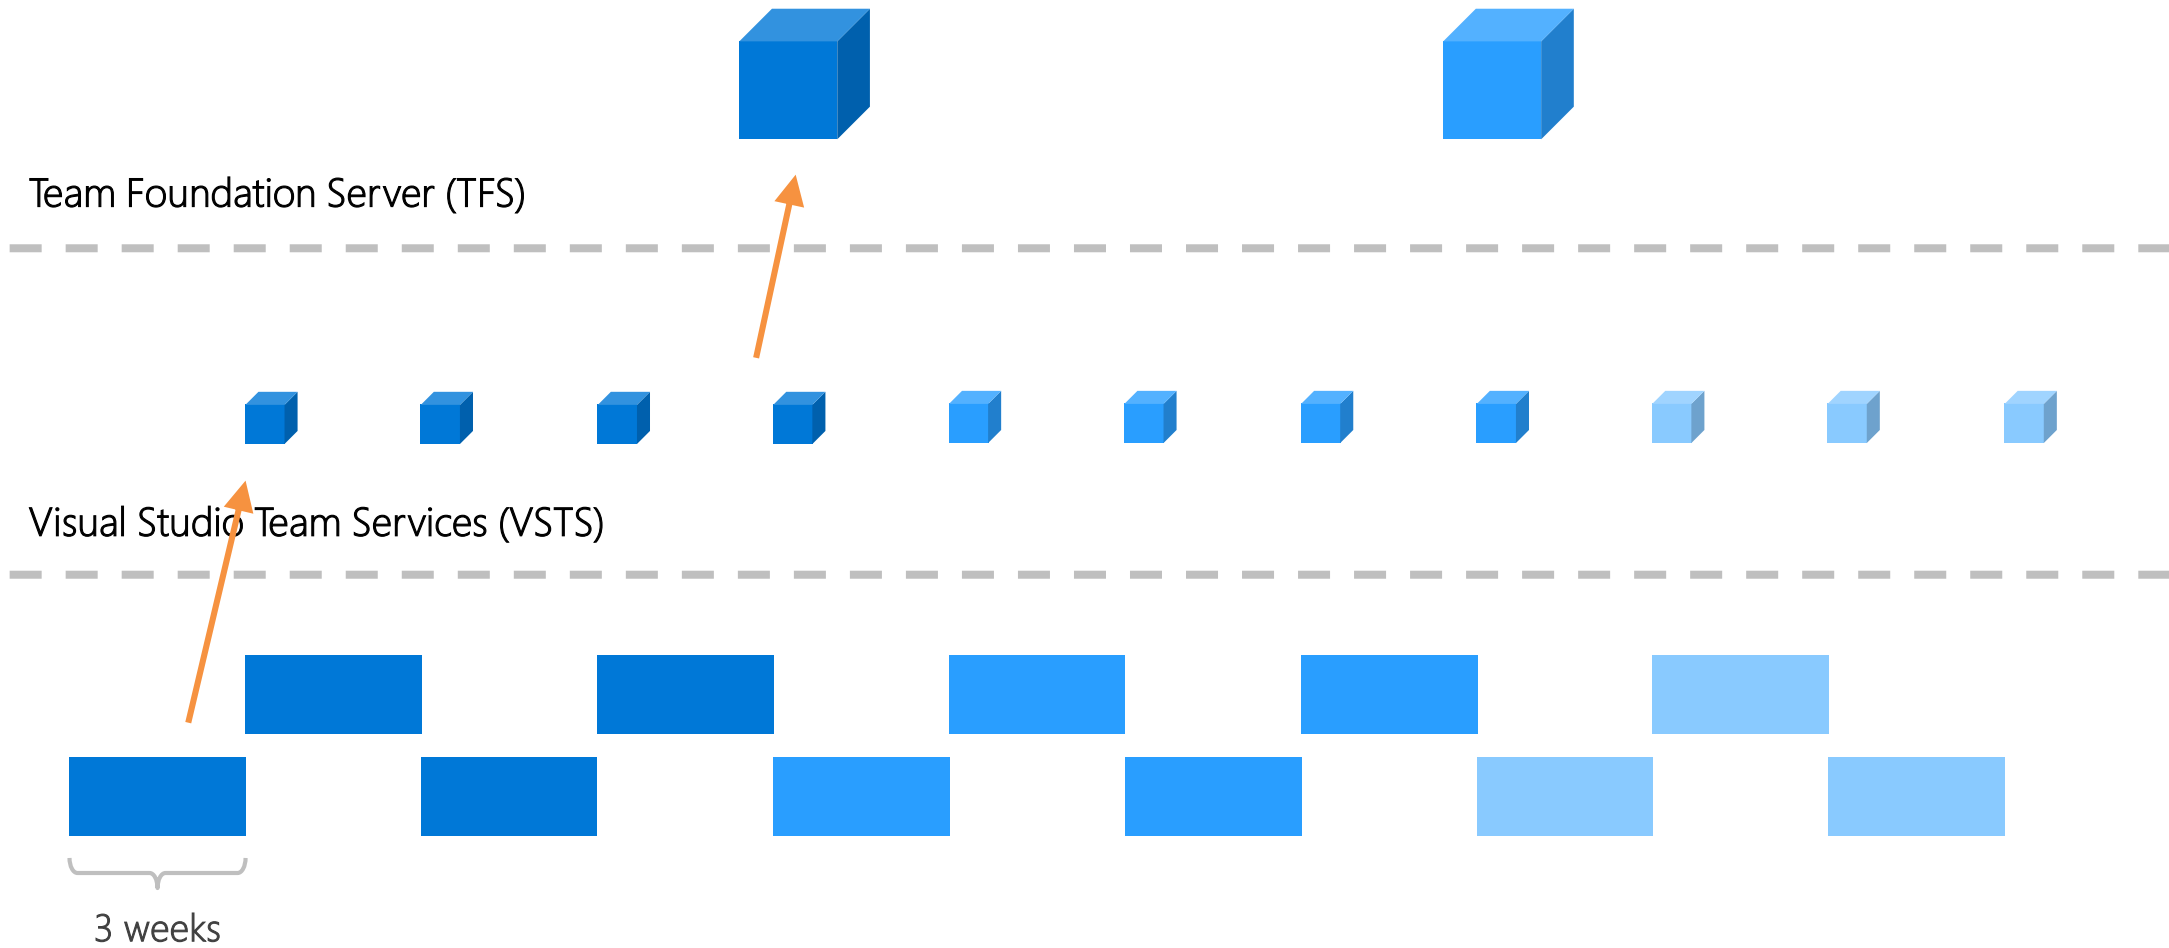
3,500

The Developer Division at Microsoft



The VSTS team... spread out across 40 feature teams

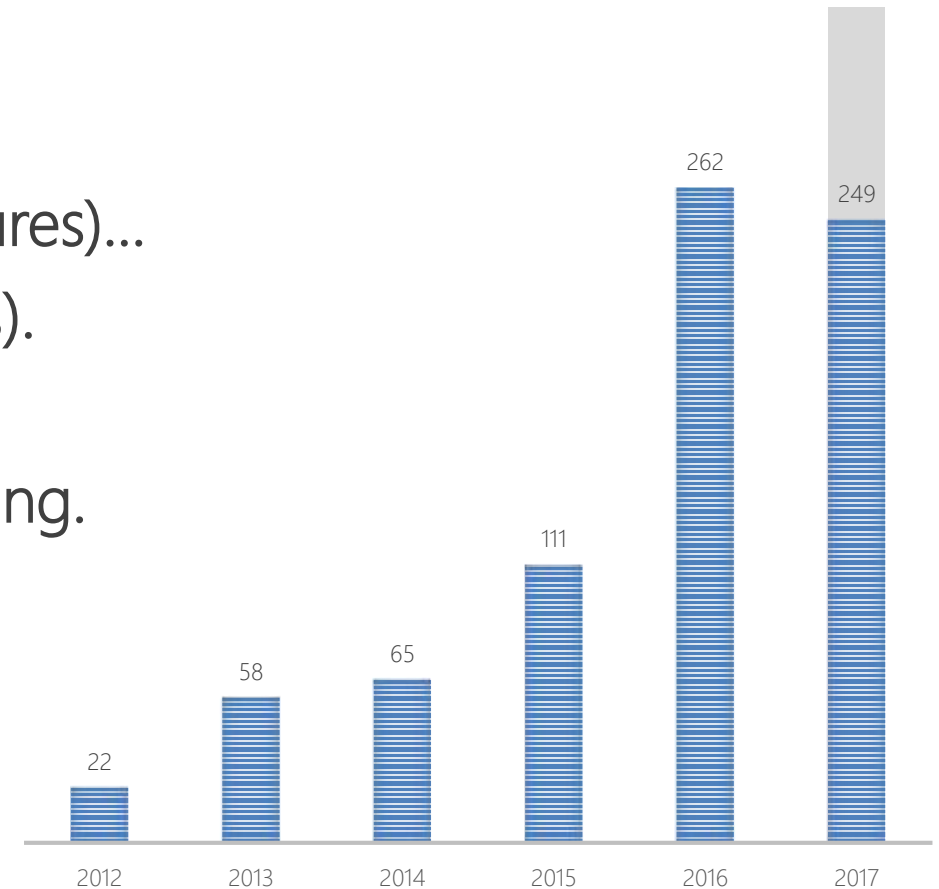
How do we work?



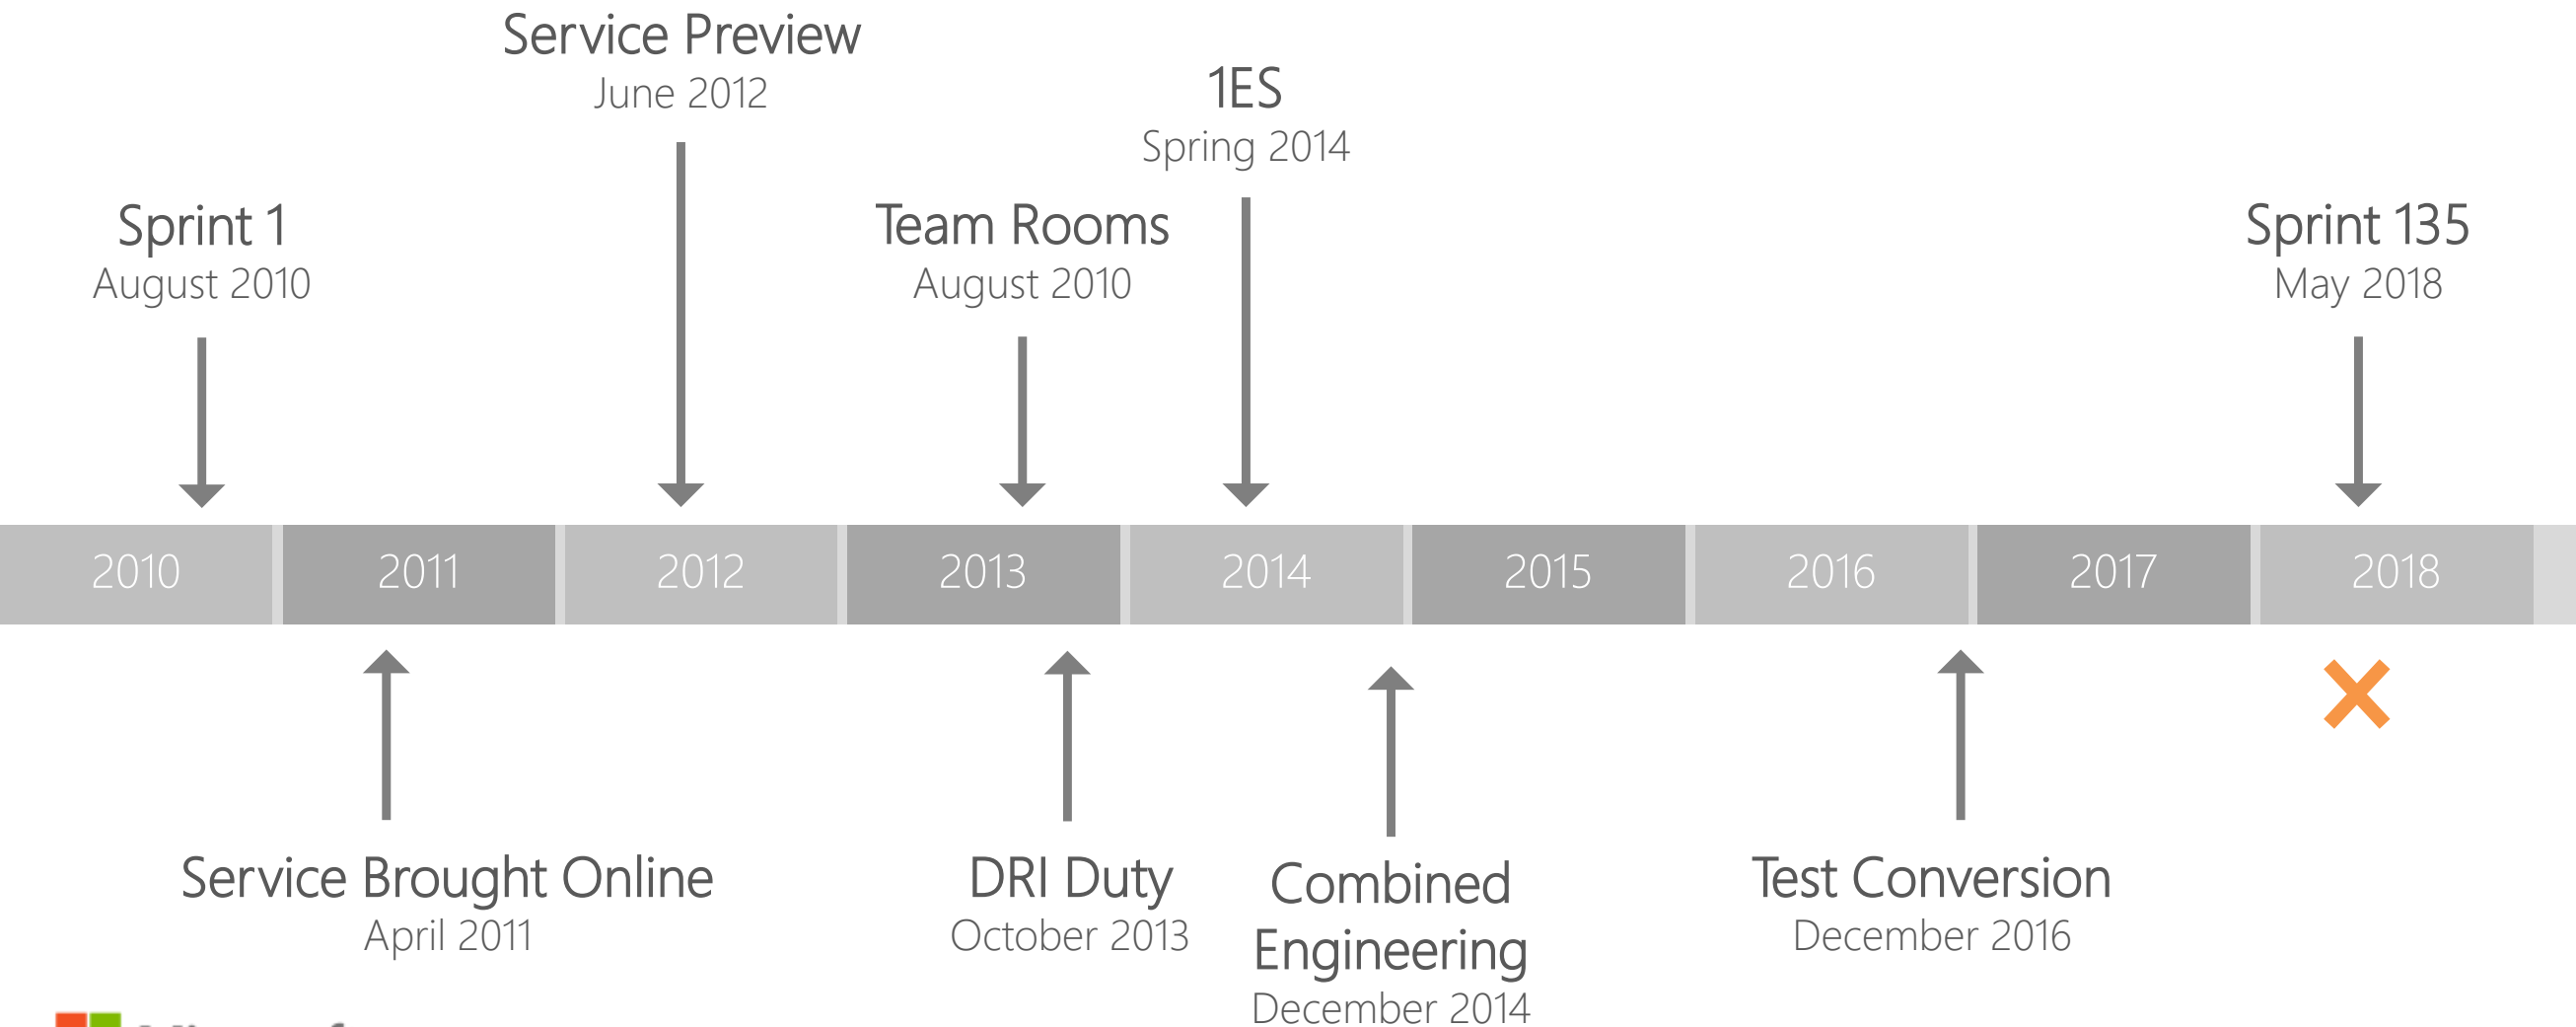
Features Delivered per Year

We are delivering value to customers and an increased velocity.

- More features in the 2016 calendar year (262 features)...
- Than the previous 4 years combined (256 features).
- 346 features in the 2017 calendar year.
- ~211 features in the 2018 calendar year and counting.



The Journey

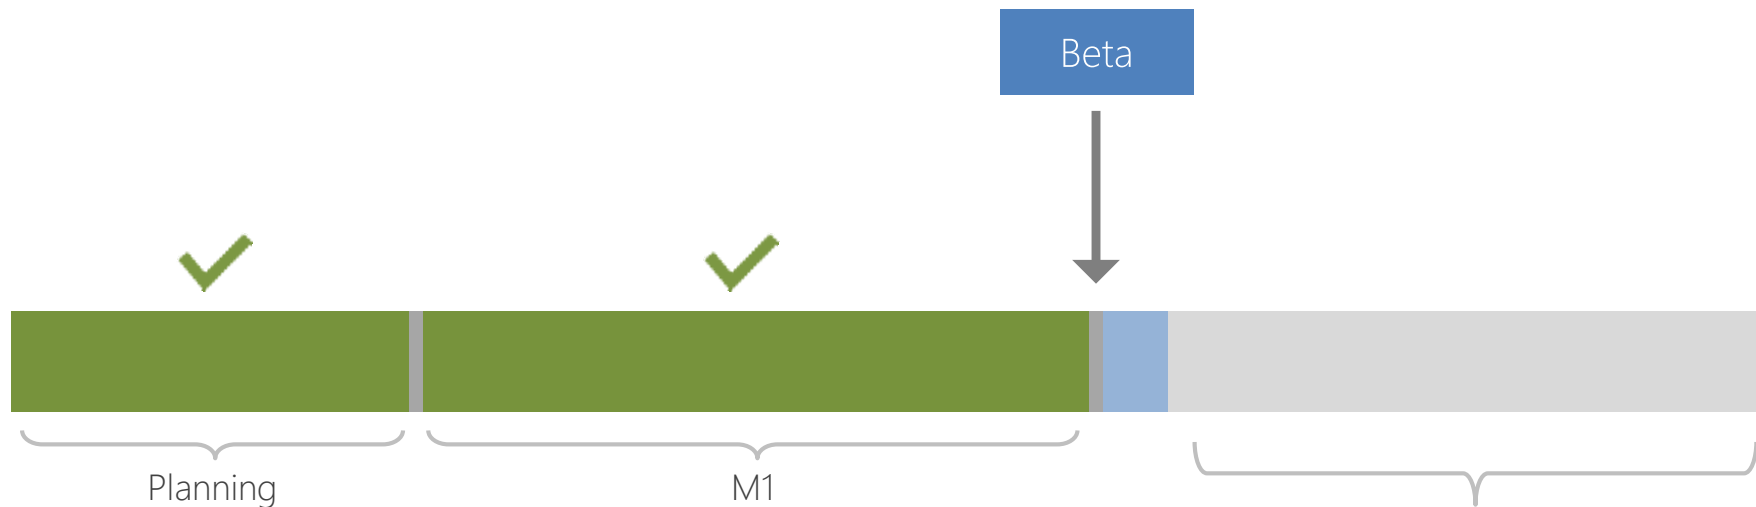


What did it look like before?

Before

The OLD way

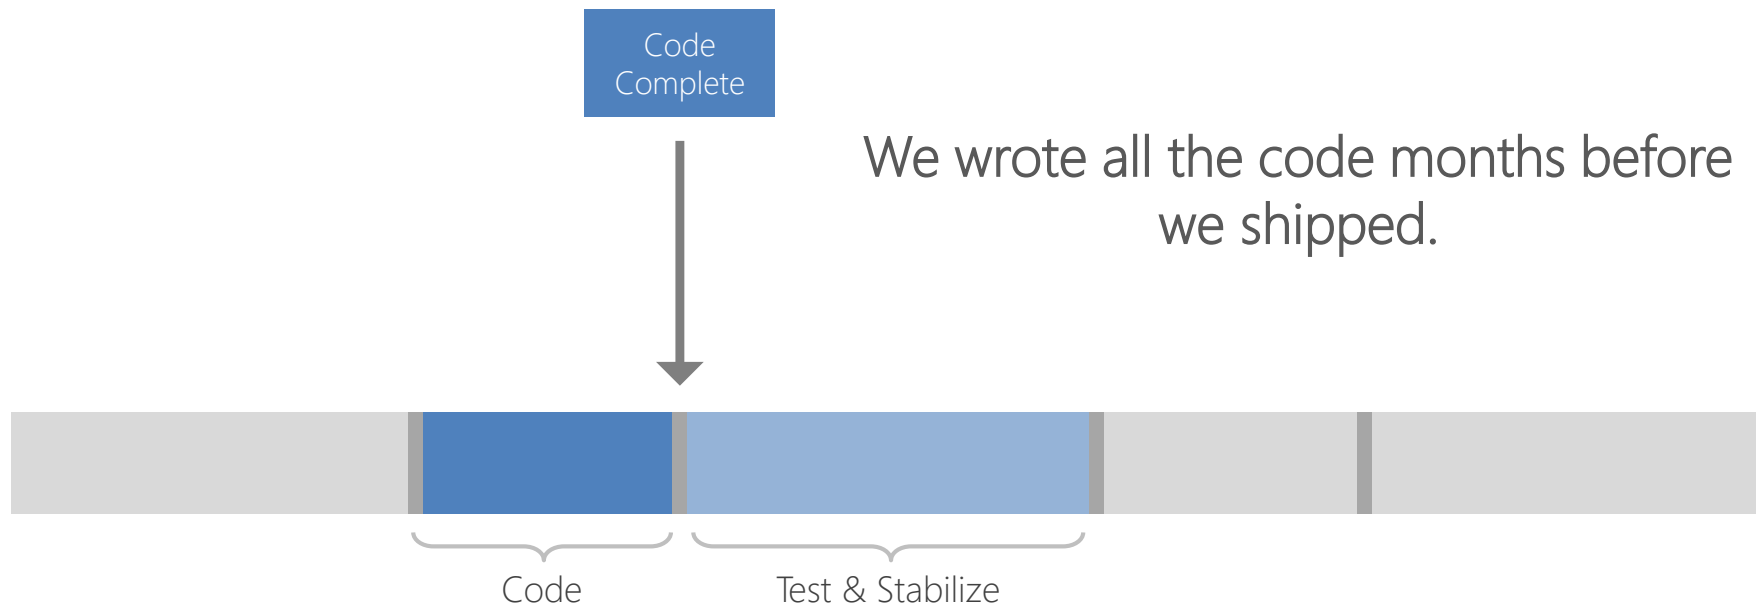
Customer feedback – we should change the way a feature works. We didn't get it *quite* right...



... but we're booked solid already.

Before

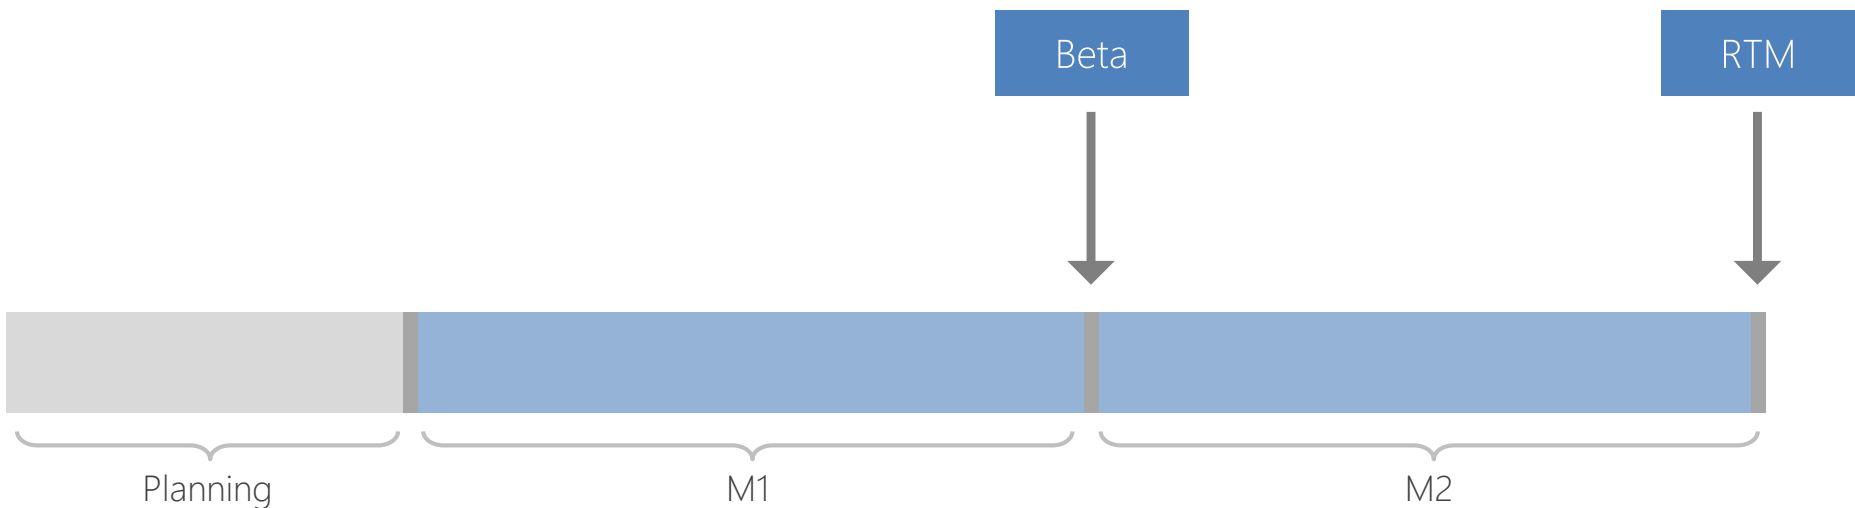
The OLD way



Before

The OLD way

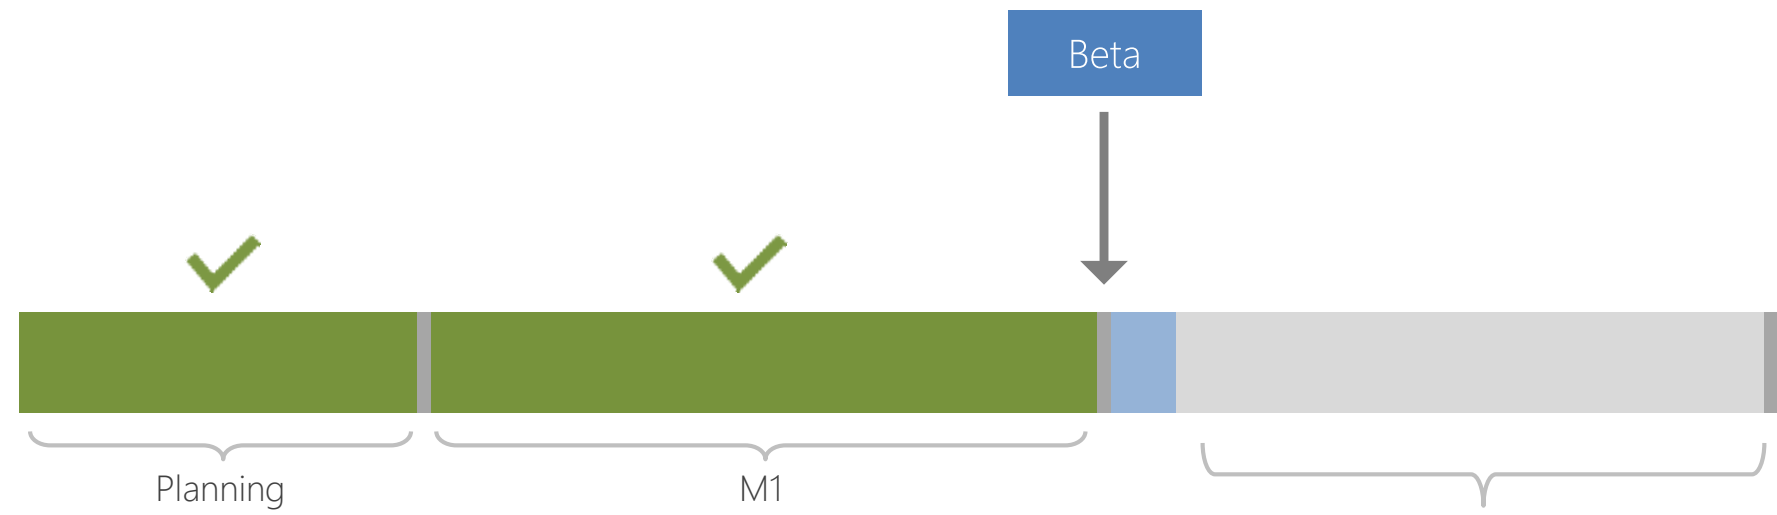
We had a perfect schedule and knew exactly when it would be ready!



Before

The OLD way

Customer feedback – we should change the way a feature works. We didn't get it *quite* right...

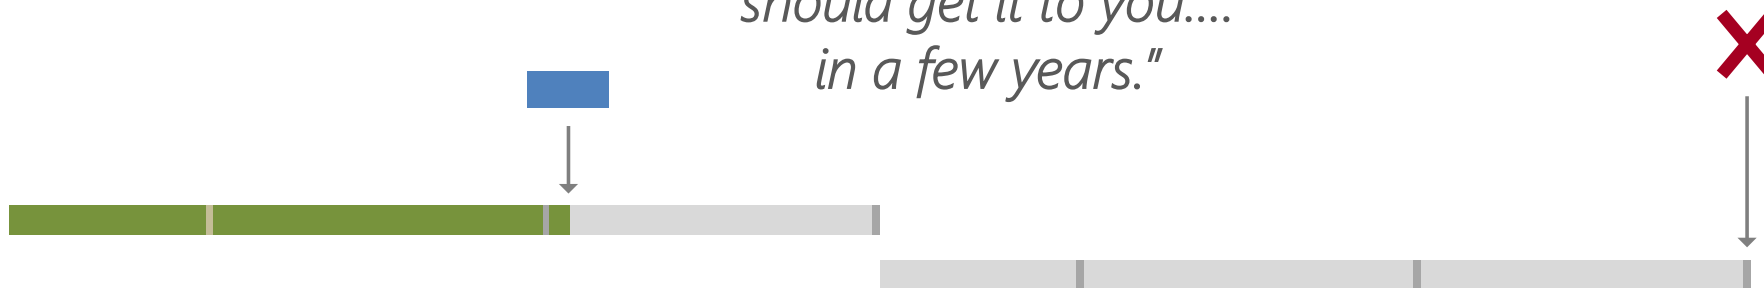


... but we're booked solid already.

Before

The OLD way

"Great feedback. Thanks! We'll take a look in planning for the next release. We should get it to you.... in a few years."



Cha Cha Cha Changes

Teams

Cross discipline

10-12 people

Self managing

Clear charter and goals

Intact for 12-18 months

Physical team rooms

Own features in production

Own deployment of features



Self Forming Teams

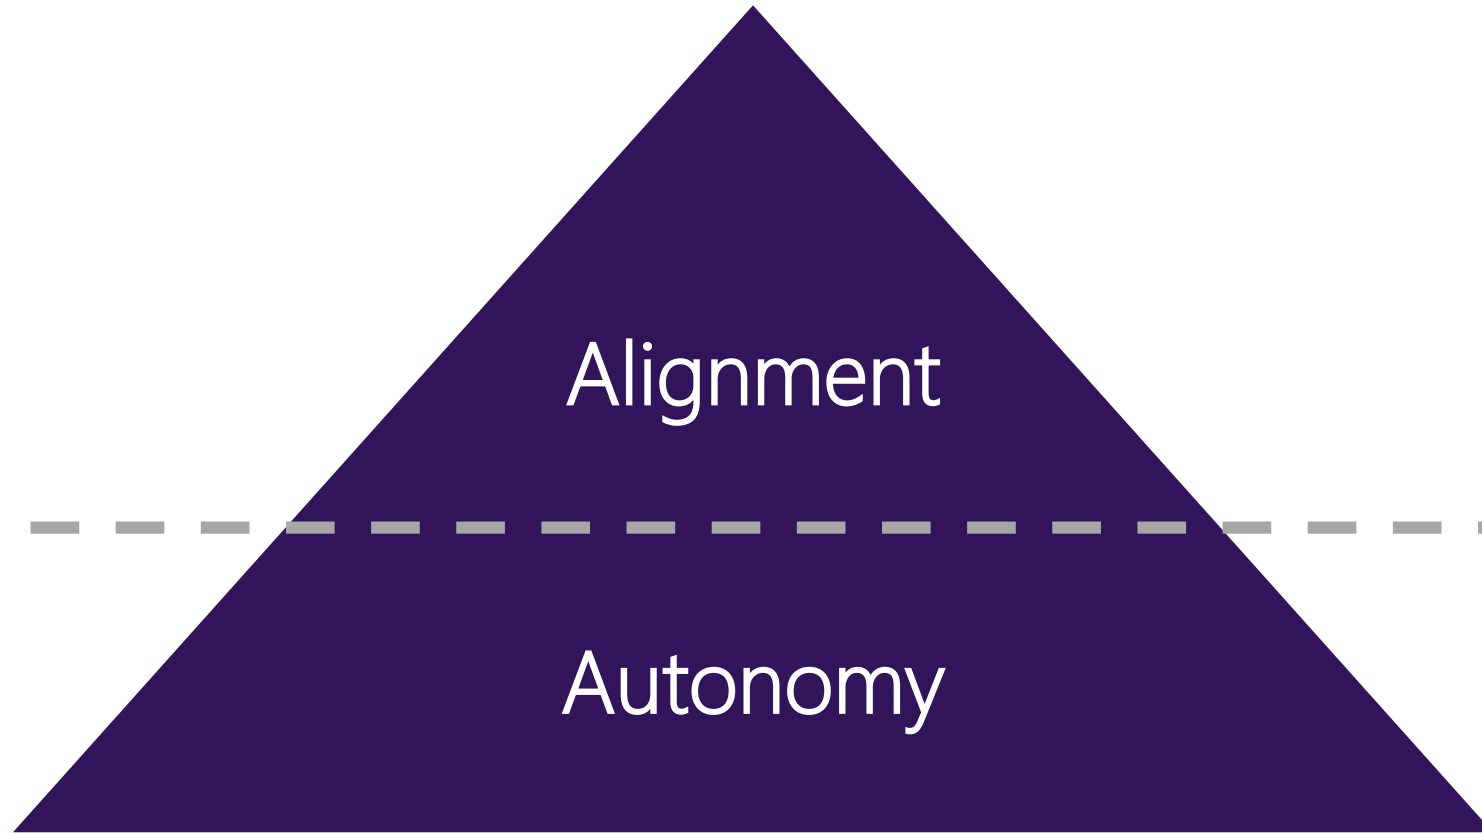
We have chosen to re-think the charter and make-up of our teams at strategic checkpoints. This happen every 18 months (or so).

The “Yellow Sticky” exercise:

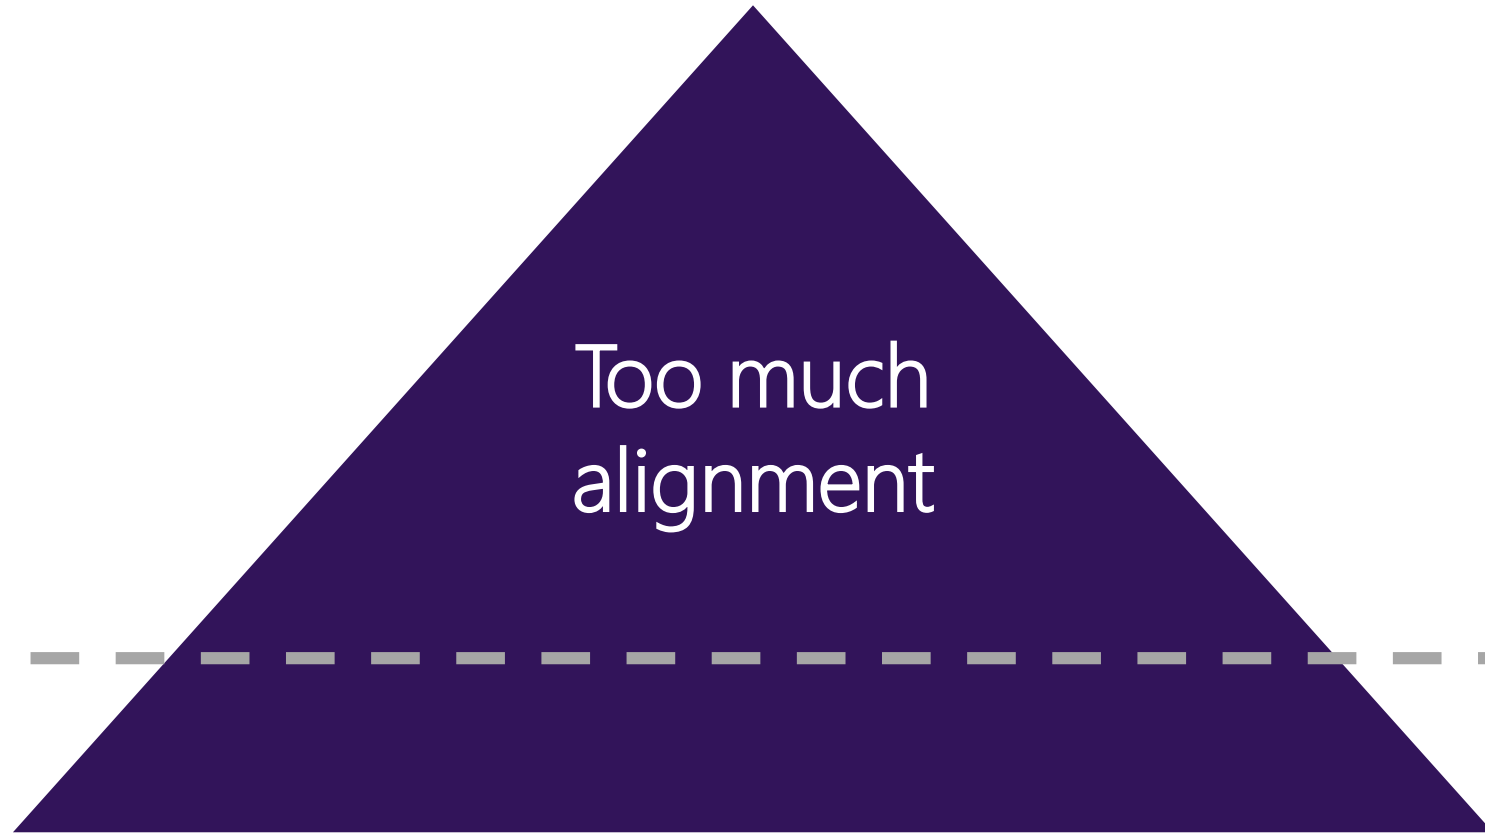
- Autonomy: Let team choose what they want to work on.
- Alignment: Ensure we’ve got the right balance across teams.



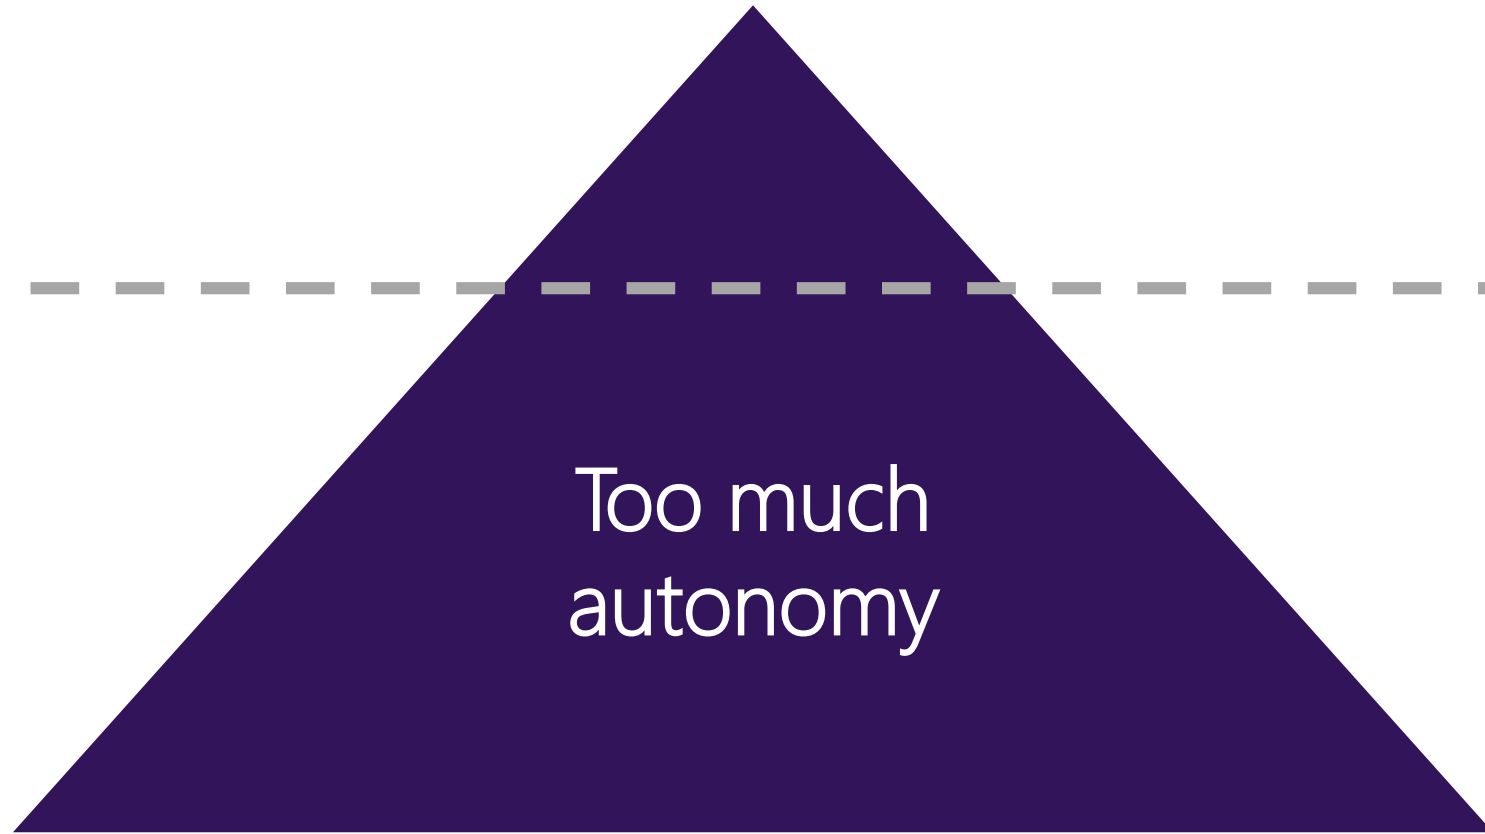
Aligned Autonomy



Aligned Autonomy




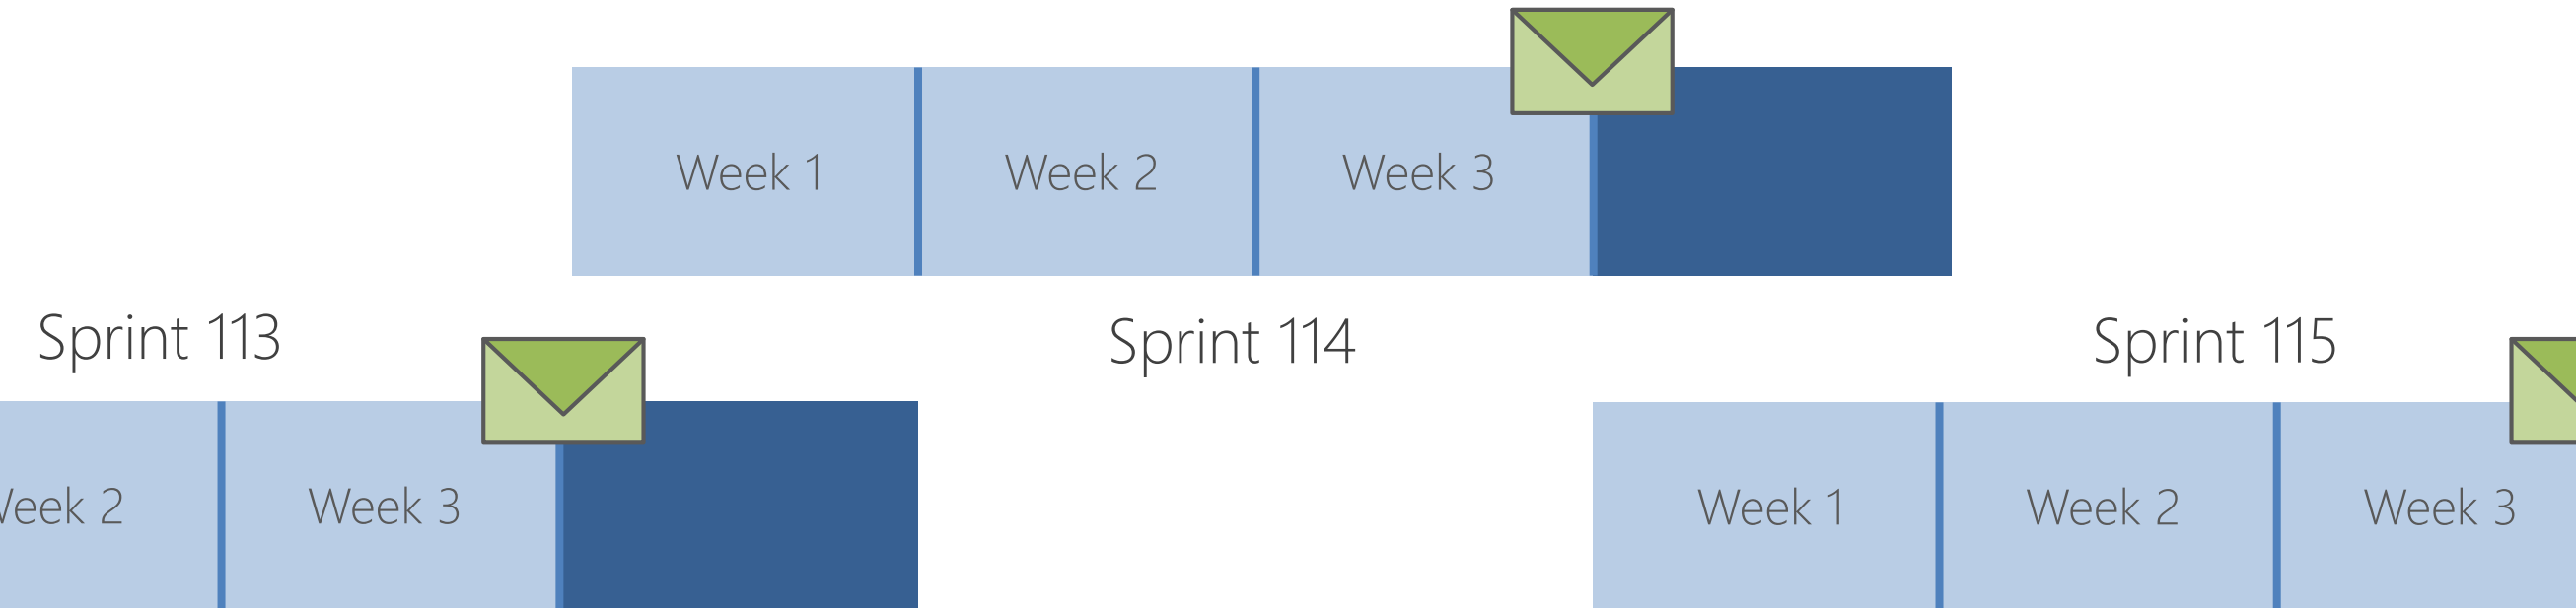
Aligned Autonomy



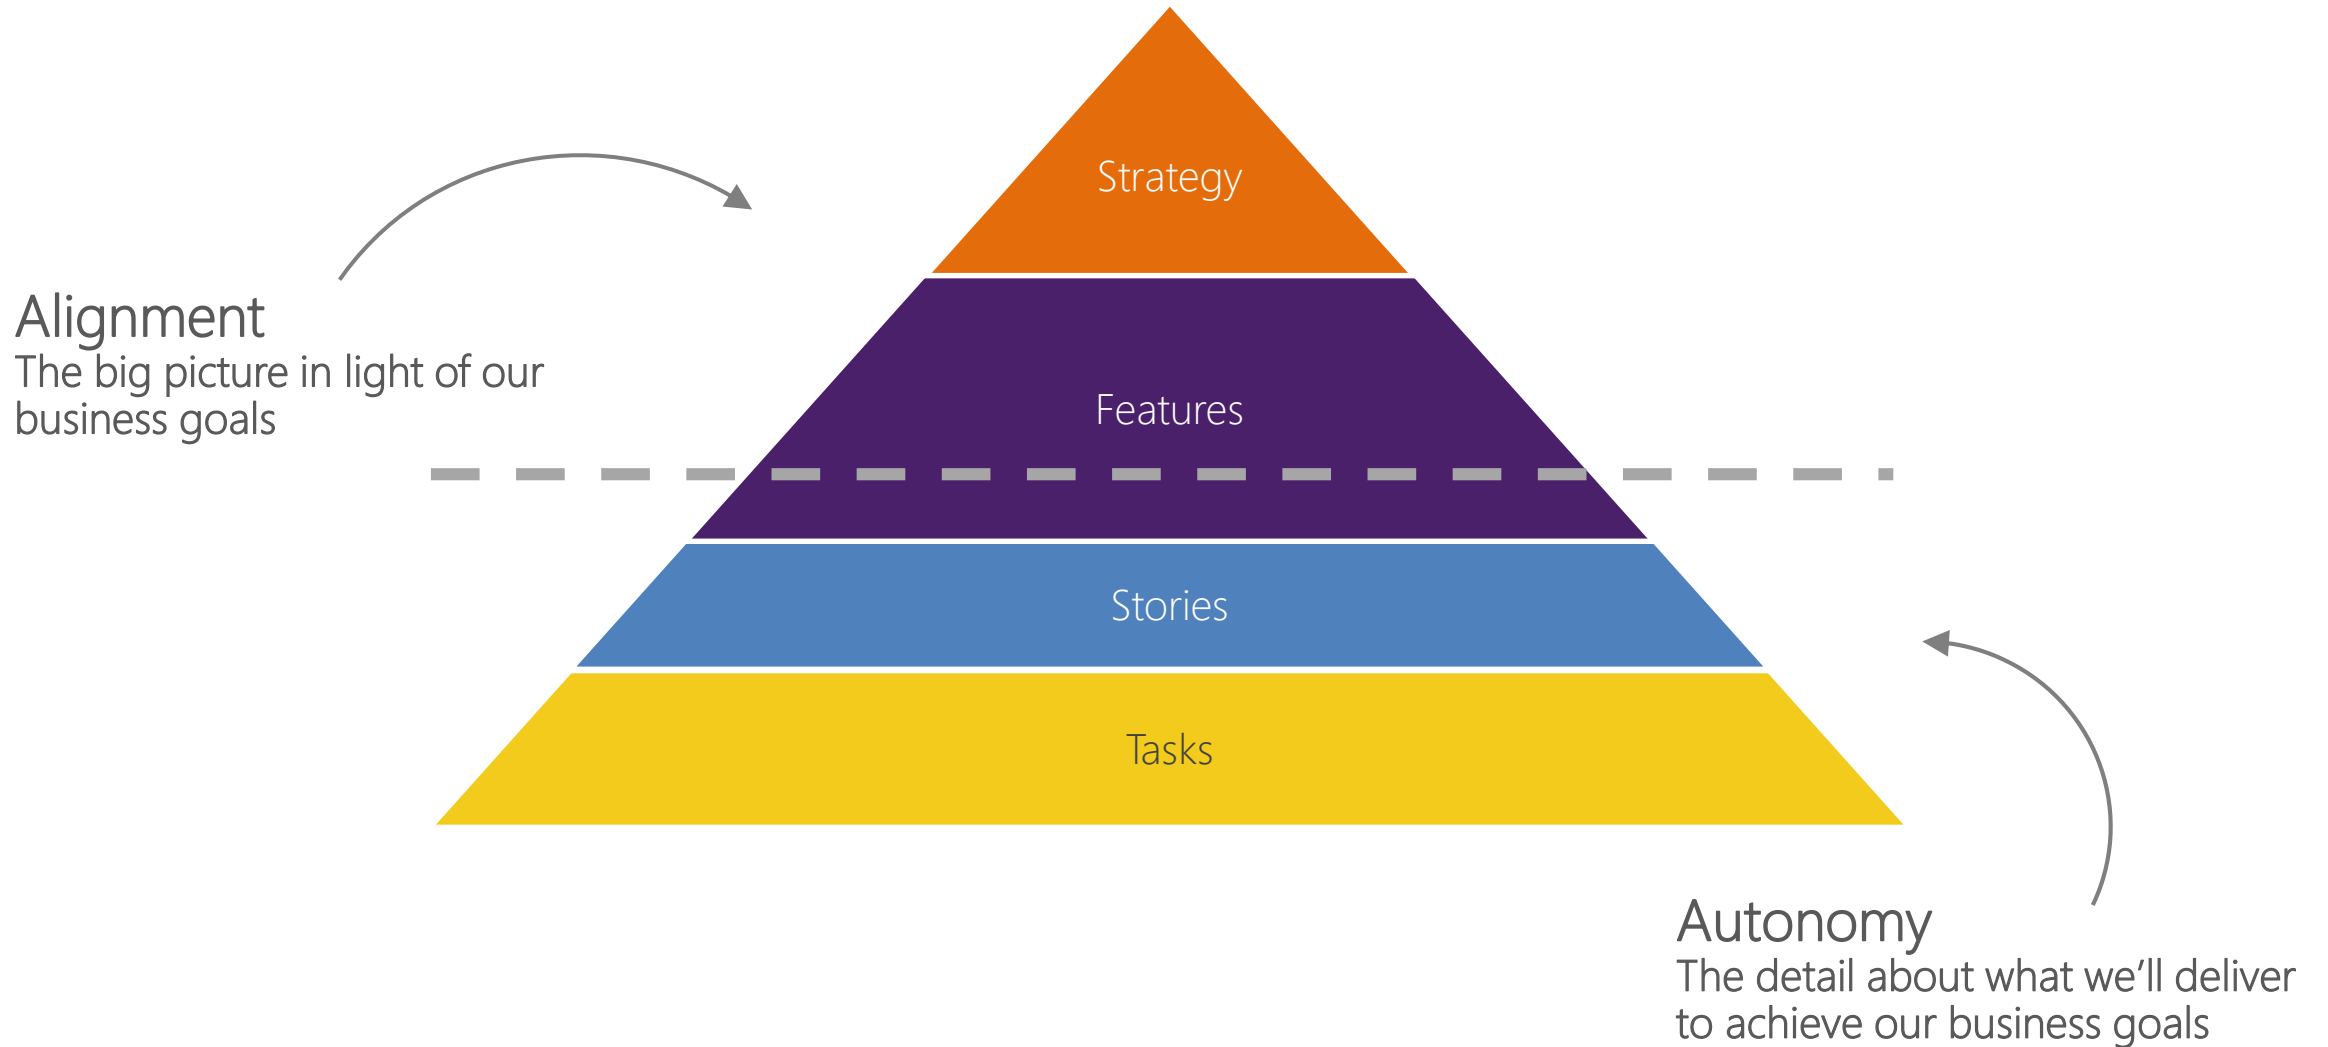
How do teams stay connected?

Sprint Mails

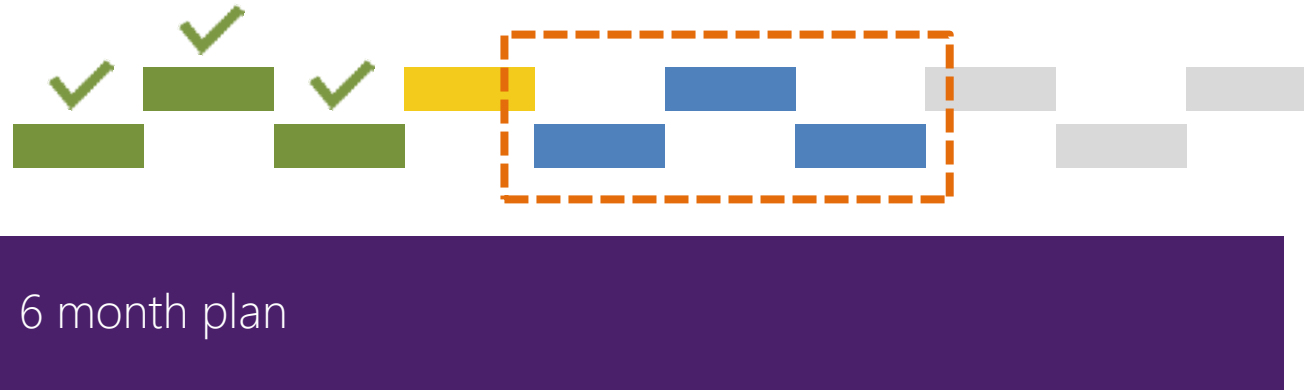
At the end of a sprint, all teams send a “sprint mail” ... communicating what they’ve accomplished in the sprint, and what they’re planning to accomplish in the next sprint.



Aligned Autonomy



Quarterly Feature Team Chats



Each team comes in and reviews with leadership three things:

1. What is the plan for the next 3-sprints?
2. Is the team healthy?
3. Any risks or issues to highlight?

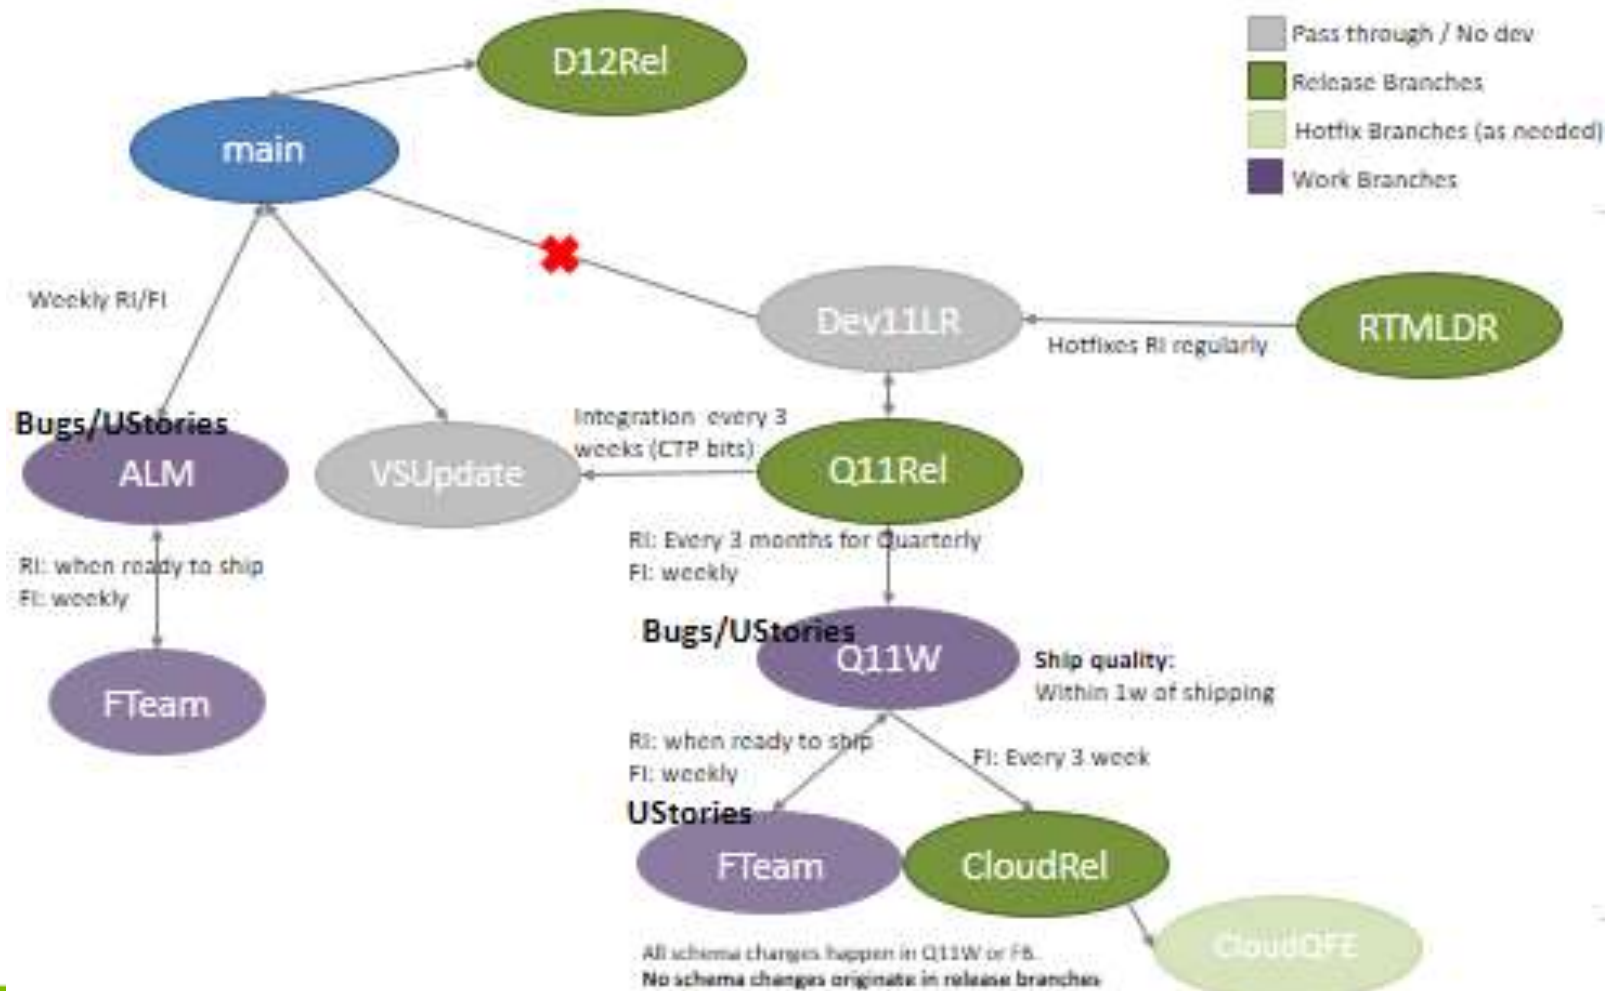
Bug Cap

We all follow a simple rule we call the “Bug Cap”:

$$\# \text{ engineers on your team} \times 5 = ?$$

Traditional branch structure

The OLD way



Deep branch hierarchy

Creates merge and integration debt

Significant costs to code flow

Complex logistics Engineers must understand

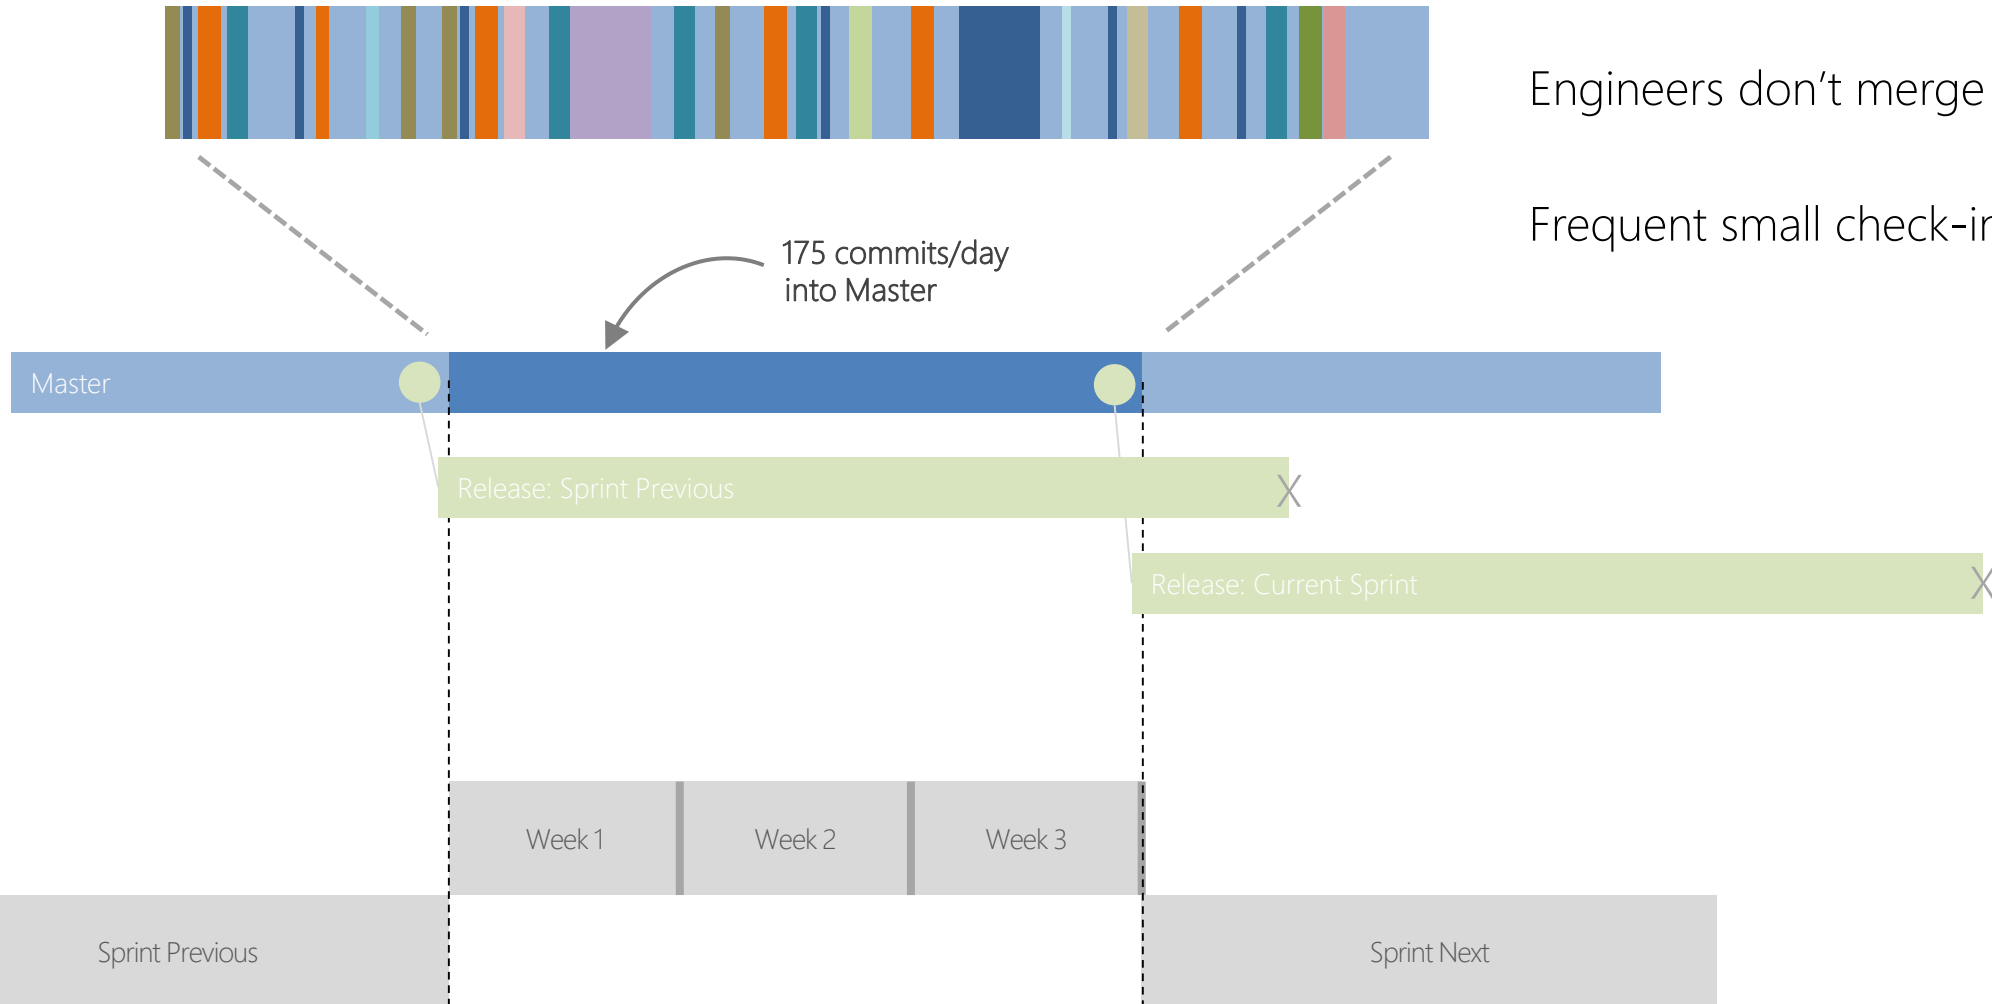
Work out of master

Source in a single git repo

All code flows to master

Engineers don't merge code they didn't write

Frequent small check-ins



Testing circa 2010 – arrival of the Cloud Services

New constraints and requirements

Faster cadence, even faster cadence, and more

Lack of customer validation through Beta, RC etc.

Micro-services deployed independently

High availability, no downtime deployments

Initial response and approach

Do the traditional waterfall dev/test model but faster

Pushed for faster automation

Test Selection techniques as a way of survival

Testing in Cloud cadence – problems

New problems emerged, old ones exacerbated

Testing became major bottleneck – we reached a breaking point

Trains didn't run on time

Lack of accountability on the Developers – no real incentive to change

High frustration among SDETs. Major retention issues.

Our model was broken

Bing, being first major cloud service at Microsoft, noticed it first

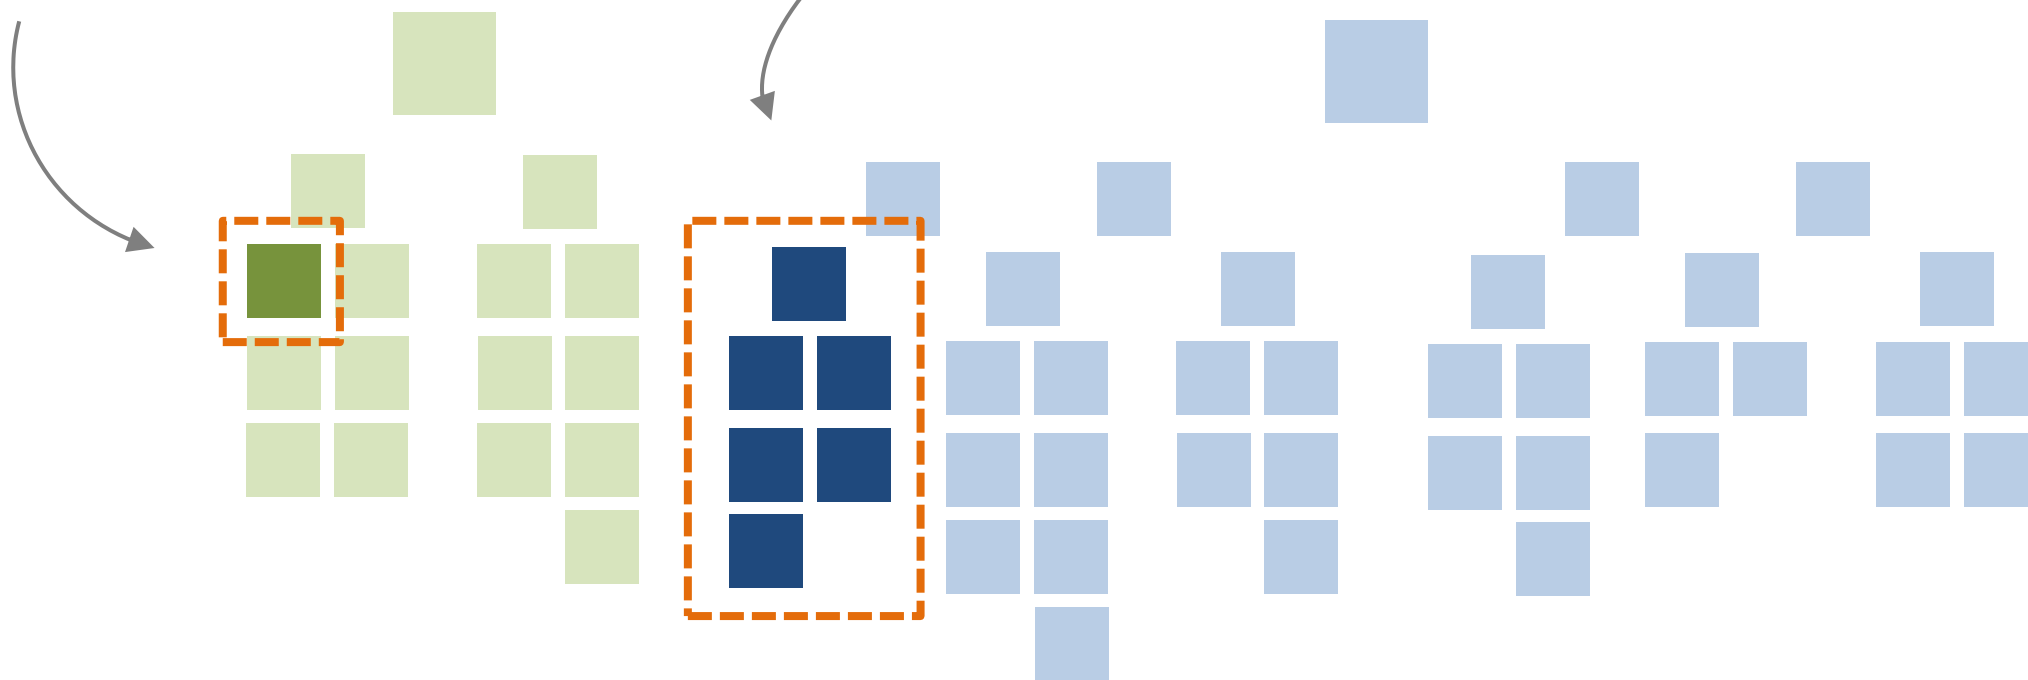
Over next few years, every team at Microsoft moved to the Cloud and changed their testing approach

Quality ownership

Teams

Program Management is responsible for:
WHAT we're building, and
WHY we're building it

Engineering is responsible for
HOW we're building it, and that
we're building it with QUALITY



Shift-Left

Master is
always
shippable

Our problems: Sept '14

Tests took too long

Over 22 hours for nightly run and 2 days for the full run

Tests failed frequently

Only ~60% of P0 runs passed 100%; Each NAR suite had many failures

Quality signal unreliable in Master

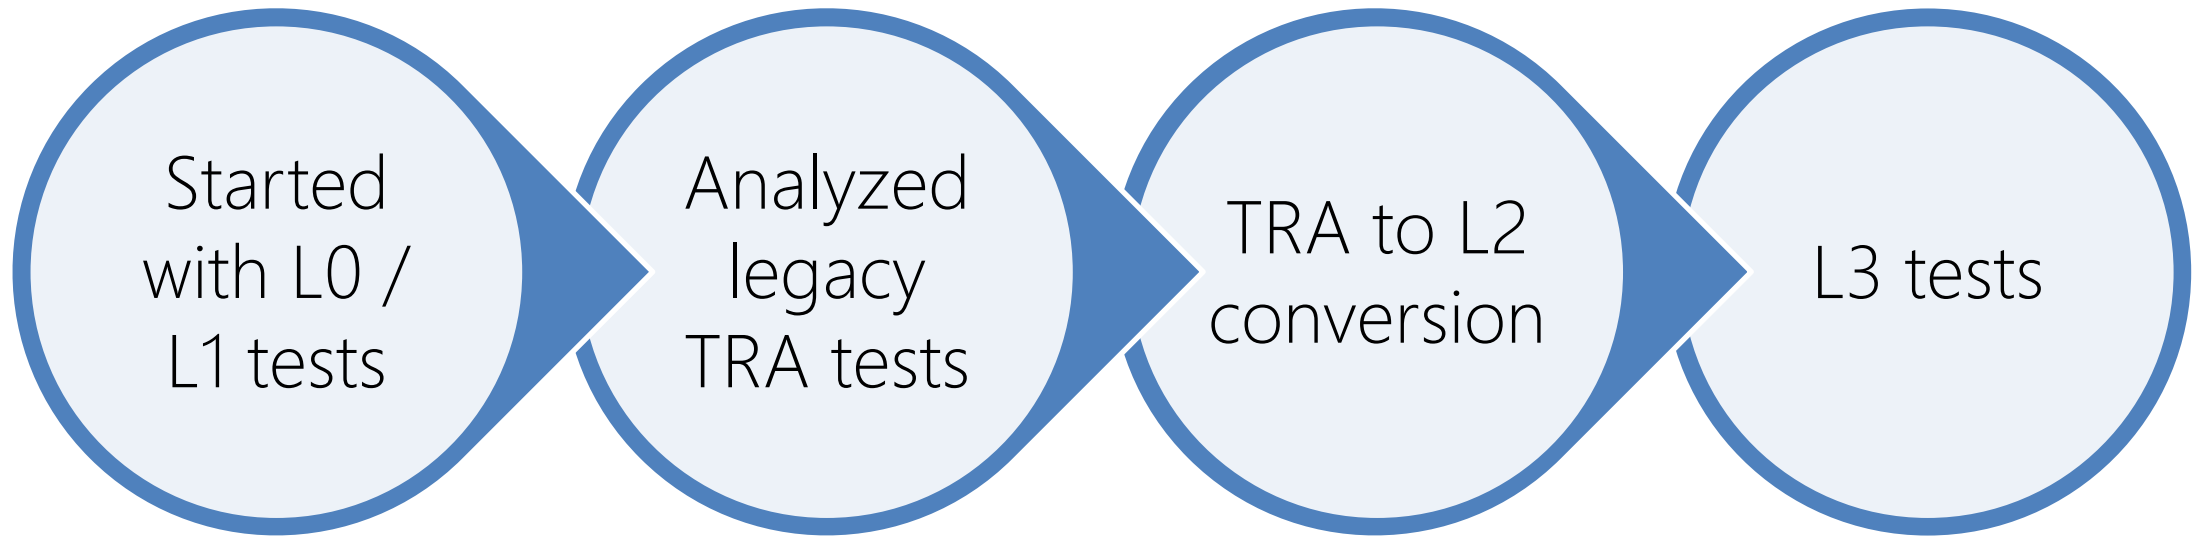
Test failure analysis was too costly

Huge cliff to reach Release quality at the end of Sprint

Took days to sift through failures before deployment could start



Shifting the test portfolio journey



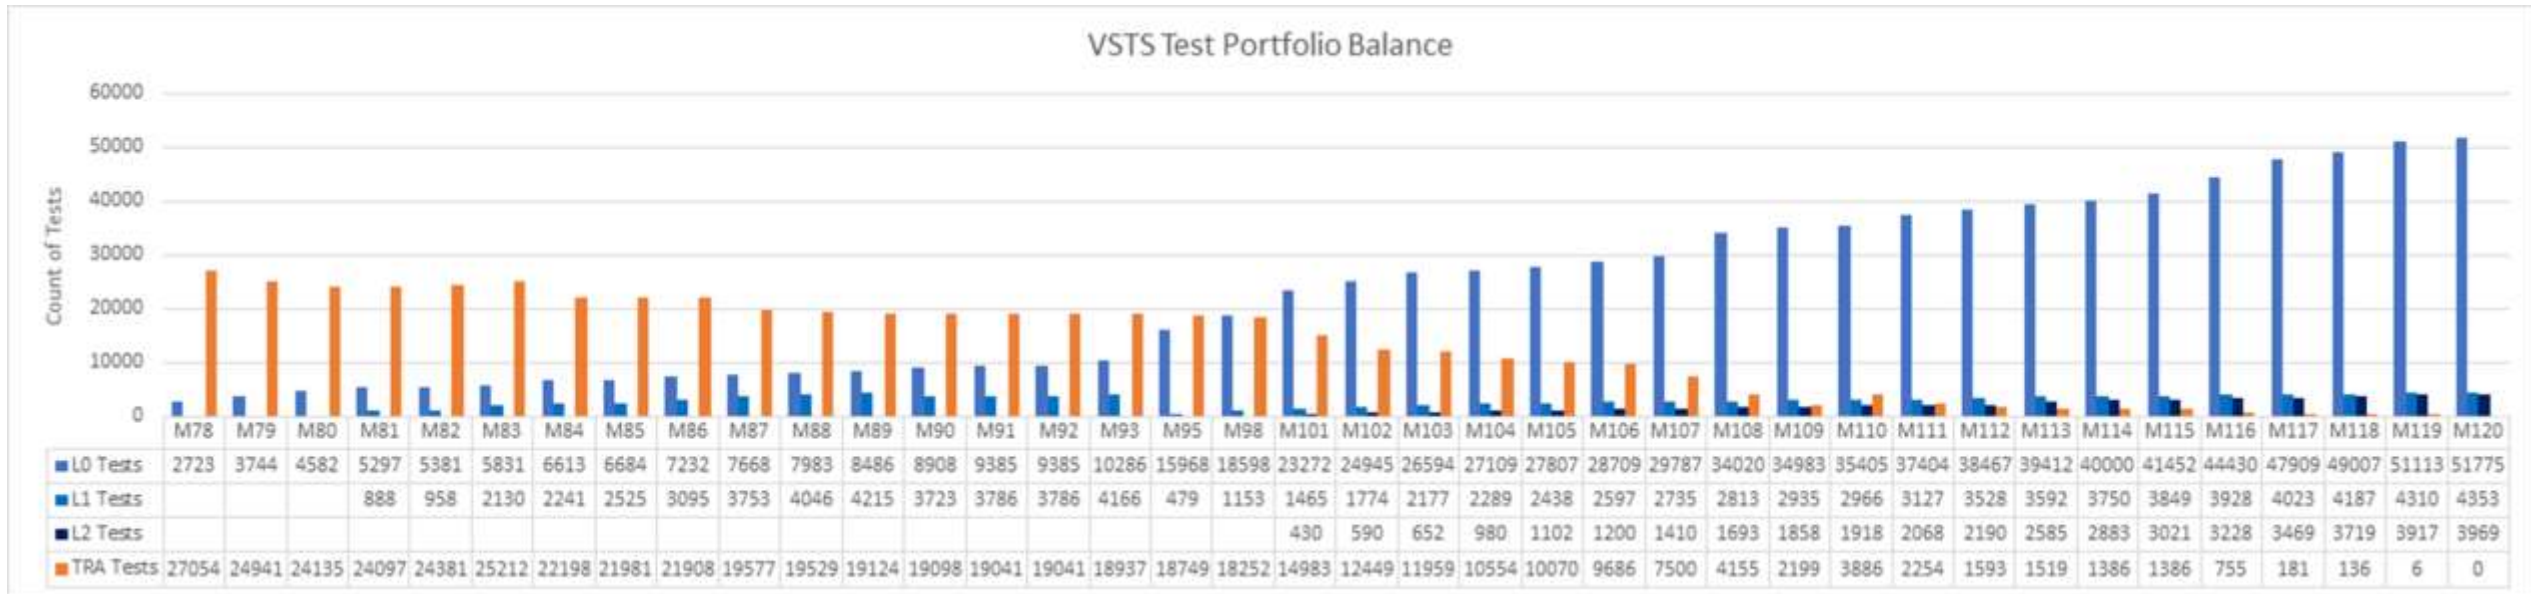
- Made it easy to author and execute high quality L0/L1 tests
- Stopped creating new TRA tests as much as possible

- Tests that can be deleted
- Tests that can move to L0/L1
- Tests that will move to VSSF Test SDK
- On-prem tests we expect to maintain in lights-on mode

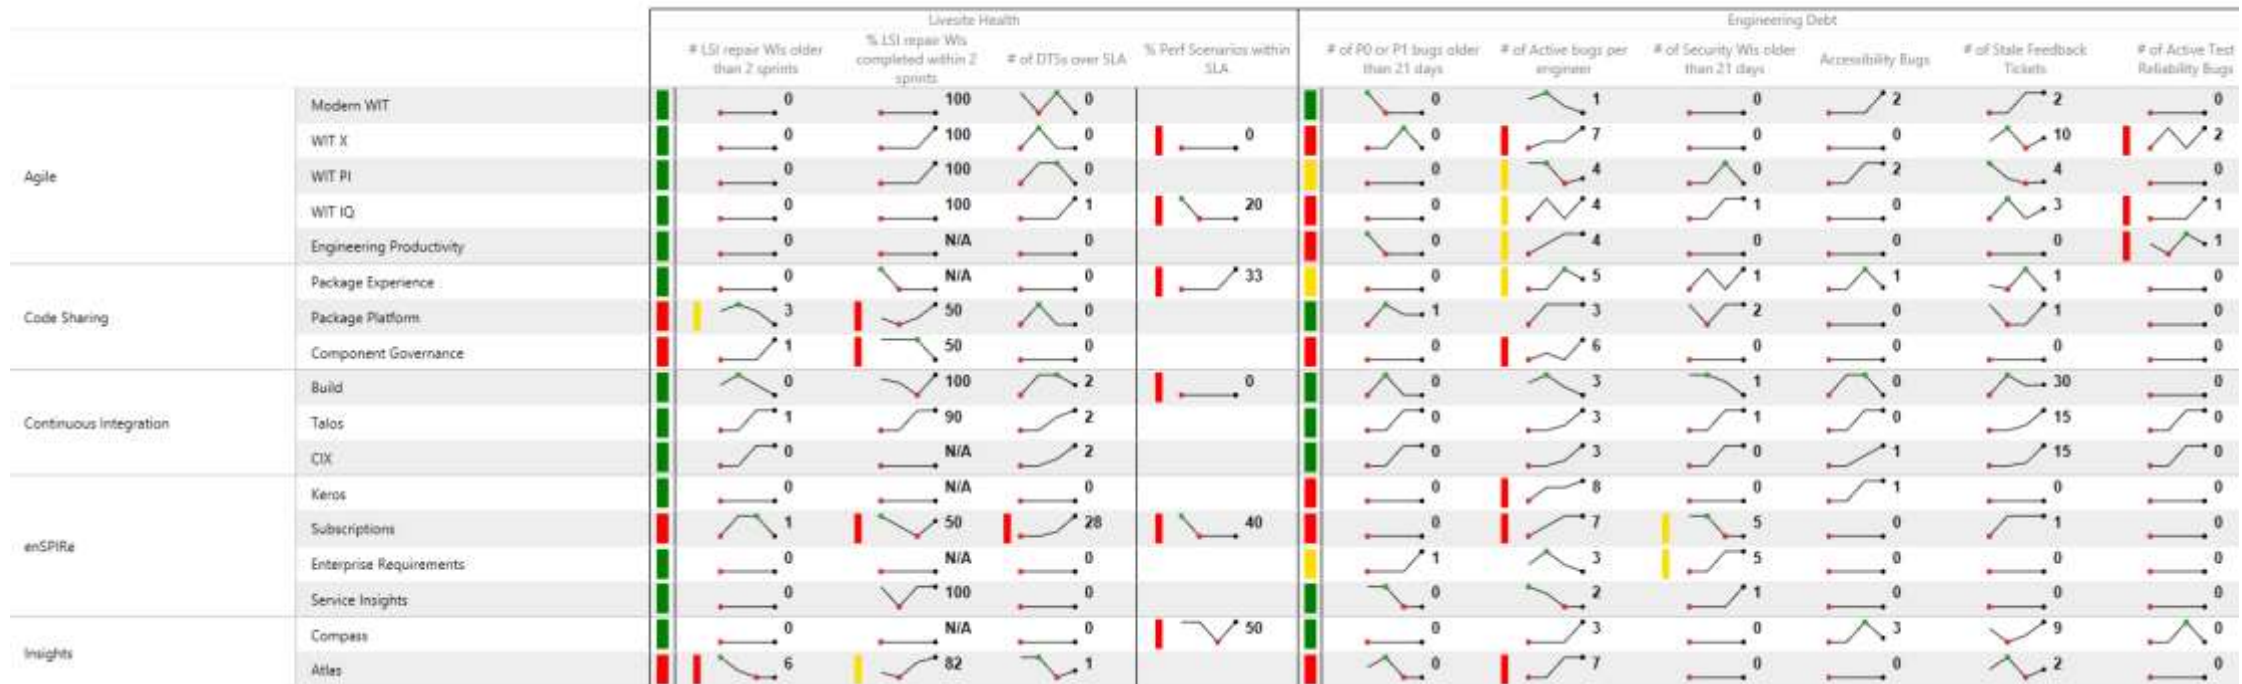
- Test Arch v-team re-wrote L2 test framework
- Top-down push from management with org wide scorecard

- Building these now

Test portfolio over time



What we track



Live Site Health/Debt

Time to Detect, Time To Mitigate

Incident prevention items

Aging live site problems

Customer support metrics (SLA, MPI, top drivers)

Engineering Health/Debt

Bug cap per engineer

Aging bugs in important categories

Pass rate & coverage

Velocity

Time to build

Time to self test

Time to deploy

Time to learn (Telemetry pipe)

Feature Flags

Controlling exposure in the cloud



What do feature flags give us?

Decouple deployment and exposure

Flags provide runtime control down to individual user

Change without redeployment

Controlled via PowerShell or web UI

Support early feedback, experimentation

Quick off switch

Control

PowerShell

Get-FeatureFlag

Set-FeatureFlag

Web UI

Tracing

FeatureFlag

Identity

Account

Build And Deployment

Users

Service Type

TFS

✓

Feature Flags

SourceControl.Revert

✓

Accounts

☒ Custom List

☐ Early Adopter Stages

hallux
buckh-westeur

✓

Display Current Status

✓

Feature Flag

☒ On

☐ Off

☐ Undefined

Submit

Account Name	Revert
hallux	On
buckh-westeur	Off

Page: 1

Deployment Practices



Early Principles

- The same tools we use to deploy to production we use in dev and test environments
- The quality signals we look at to green light deployments are tracked constantly every day
- Deployments take zero down time
- Deployments happen during working hours

Are We Ready To Deploy?

Check for blocking bugs

Check test results

Choose a good build

Release Branch Runs - Default

Environments

Sps.SelfTest

Sps.SelfHost

Tfs.SelfHost Set 1

Tfs.SelfHost Set 2

Tfs.SelfTest

Tfs.Deploy

TfsOnPrem.SelfHost

TfsOnPrem.SelfTest

✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	▶
✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	▶	▶
✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	▶	▶
✓ 100%	✓ 100%	✓ 100%	✓ 100%	✗ 98.69%	✓ 100%	▶	▶
✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	▶	▶
	✓ 100%	✓ 100%	✓ 100%		✓ 100%	▶	
✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	▶	▶
✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	✓ 100%	▶	▶

Safe Deployment

What is Safe Deployment?

- Deploy changes to risk tolerant customers first, progressively roll out to larger and larger sets of customers
- Automated health checks and roll back

Deployment Rings

Ring	Purpose	Customer type	Data center
0	Surface most of the customer-impacting bugs introduced by the deployment	Internal only, high tolerance for risk and bugs	US West Central
1	Surface bugs in areas that we do not dogfood	Customers using a breadth of the product, especially areas we do not dogfood (TFVC, hosted build, etc). Should be in a US time zone.	A small data center
2	Surface scale-related issues	Public accounts. Ideally free accounts, using a diverse set of the features available.	A medium to large US data center
3	Surface scale issues common in internal accounts and international related issues	Large internal accounts European accounts	Internal data center and a European data center
4	Update the remaining scale units	Everyone else	All the rest

Wrap-Up



Before

4-6 month milestones
Horizontal teams
Personal offices
Long planning cycles
PM, Dev, Test
Yearly customer engagement
Feature branches
20+ person teams
Secret roadmap
Bug debt
100 page spec documents
Private repositories
Deep organizational hierarchy
Success is a measure of install numbers
Features shipped once a year

After

3-week sprints
Vertical teams
Team rooms
Continual Planning & Learning
PM & Engineering
Continual customer engagement
Everyone in master
8-12 person teams
Publicly shared roadmap
Zero debt
Specs in PPT
Open source
Flattened organization hierarchy
User satisfaction determines success
Features shipped every sprint

Resources

<https://aka.ms/devops>

<https://youtube.com/devopsatmicrosoft>

<https://aka.ms/devopslab>

Thank you